

Contents

cc-manage Provider Setup Guide	2
Quick Reference	3
1. Anthropic	3
Get API Key	3
Create Profile	3
Launch	4
2. Gemini	4
Get API Key	4
Create Profile	4
How It Works	4
Launch	4
3. OpenRouter	4
Get API Key	5
Create Profile	5
Notes	5
Launch	5
4. Groq	5
Get API Key	5
Create Profile	5
Notes	6
Launch	6
5. Mistral	6
Get API Key	6
Create Profile	6
Default Models	6
Launch	7
6. Mistral Vibe	7
Get API Key	7
Create Profile	7
Default Models	7
Notes	7
Launch	7
7. Codestral	7
Get API Key	8
Create Profile	8
Default Models	8
Notes	8
Launch	8
8. OpenCode Nemotron	8
Get Auth Token	8
Create Profile	8
Default Model	9
How It Works	9
Launch	9
9. NVIDIA NIM	9
Get API Key	9

Create Profile	9
Default Models	9
Launch	10
10. Hugging Face	10
Get API Key	10
Create Profile	10
Default Models	10
Notes	10
Launch	10
11. DeepSeek	10
Get API Key	11
Create Profile	11
Default Models	11
Launch	11
12. Fireworks	11
Get API Key	11
Create Profile	11
Default Model	12
Launch	12
13. Together	12
Get API Key	12
Create Profile	12
Default Model	12
Launch	12
14. xAI	12
Get API Key	12
Create Profile	13
Default Models	13
Launch	13
15. Ollama Cloud	13
Get API Key	13
Create Profile	13
Default Models	13
Launch	13
16. Custom OpenAI-Compatible	14
Create Profile	14
Requirements	14
Launch	14
Provider Comparison	14
Response Format Support	14
Speed Tiers	15

cc-manage Provider Setup Guide

Step-by-step instructions for configuring each supported LLM provider with cc-manage.

Quick Reference

Provider	API Key Source	Mode	Proxy
Anthropic	console.anthropic.com	Direct	No
Gemini	aistudio.google.com	Proxy	Yes
OpenRouter	openrouter.ai/keys	Direct	No
Groq	console.groq.com	Proxy	Yes
Mistral	console.mistral.ai	Proxy	Yes
Mistral Vibe	console.mistral.ai	Proxy	Yes
Codestral	console.mistral.ai	Proxy	Yes
OpenCode Nemotron	opencode.ai	Proxy	Yes
NVIDIA NIM	build.nvidia.com	Proxy	Yes
Hugging Face	huggingface.co/settings/tokens	Proxy	Yes
DeepSeek	platform.deepseek.com	Direct	No
Fireworks	fireworks.ai	Direct	No
Together	api.together.xyz	Proxy	Yes
xAI	console.x.ai	Proxy	Yes
Ollama Cloud	ollama.com	Proxy	Yes
Custom	Any OpenAI-compatible endpoint	Proxy	Yes
OpenAI-compatible			

1. Anthropic

The simplest setup — no proxy needed.

Get API Key

1. Go to console.anthropic.com
2. Navigate to **API Keys**
3. Click **Create Key**
4. Copy the key (starts with sk-ant-)

Create Profile

```
cc-manage add
```

1. Select **Anthropic** from the provider list
2. Enter filename (e.g., anthropic-main)
3. When prompted for API key, choose "Add a new API key"
4. Enter a suffix (e.g., main) and paste your API key
5. Accept the default base URL: <https://api.anthropic.com>
6. Select models (defaults: claude-sonnet-4-20250514, claude-opus-4-20250514)
7. Choose default model

Launch

```
cc-switch anthropic-main  
ccs
```

2. Gemini

Uses a local proxy to translate Anthropic Messages to Gemini generateContent.

Get API Key

1. Go to aistudio.google.com/apikey
2. Click **Create API Key**
3. Copy the key (starts with AIza)

Create Profile

```
cc-manage add
```

1. Select **Gemini** from the provider list
2. Enter filename (e.g., `gemini-flash`)
3. Add new API key with suffix (e.g., `google`)
4. Accept the default base URL: `http://127.0.0.1:18000` (local proxy)
5. Select models (defaults: `gemini-2.5-flash`, `gemini-3.5-flash`)
6. Choose default model

How It Works

- `cc-manage` starts a local proxy on port 18000
- The proxy translates Anthropic Messages → Gemini generateContent format
- Supports: text, images, tool calls, tool results, streaming
- API key is sent via `x-goog-api-key` header (not URL)

Launch

```
cc-switch gemini-flash  
ccs
```

3. OpenRouter

Uses OpenRouter's Anthropic-compatible endpoint directly.

Get API Key

1. Go to openrouter.ai/keys
2. Click **Create Key**
3. Copy the key (starts with sk-or-)

Create Profile

```
cc-manage add
```

1. Select **OpenRouter** from the provider list
2. Enter filename (e.g., openrouter-free)
3. Add new API key with suffix (e.g., or)
4. Accept the default base URL: <https://openrouter.ai/api>
5. Select models (defaults: anthropic/claude-sonnet-4, google/gemini-2.5-flash)
6. Choose default model

Notes

- OpenRouter's Anthropic-compatible path works directly — no proxy needed
- Free models available with :free suffix (e.g., google/gemini-2.0-flash:free)
- Use `cc-manage models openrouter --refresh` to see all available models

Launch

```
cc-switch openrouter-free  
ccs
```

4. Groq

Fast inference with dynamic model discovery.

Get API Key

1. Go to console.groq.com/keys
2. Click **Create API Key**
3. Copy the key (starts with gsk_)

Create Profile

```
cc-manage add
```

1. Select **Groq** from the provider list
2. Enter filename (e.g., groq-fast)
3. Add new API key with suffix (e.g., groq)
4. Accept the default base URL: <http://127.0.0.1:18100> (local proxy)

5. Choose whether to fetch live models: y
6. Select from available models
7. Choose default model

Notes

- Groq uses OpenAI-compatible API through the local proxy
- Output tokens are capped at 4096 by default (Groq limitation)
- Oversized requests (>32MB) are rejected locally before hitting the API
- Use `cc-manage models groq --refresh` to update the model list

Launch

```
cc-switch groq-fast  
ccs
```

5. Mistral

Mistral's full model lineup through local proxy.

Get API Key

1. Go to console.mistral.ai
2. Navigate to **API Keys**
3. Create a new key

Create Profile

```
cc-manage add
```

1. Select **Mistral** from the provider list
2. Enter filename (e.g., `mistral-large`)
3. Add new API key with suffix
4. Accept the default base URL: `http://127.0.0.1:18005` (local proxy)
5. Fetch live models or use defaults
6. Choose default model

Default Models

- `mistral-large-latest`
- `pixtral-large-latest`
- `ministral-8b-latest`

Launch

```
cc-switch mistral-large  
ccs
```

6. Mistral Vibe

Mistral's Vibe model family with dedicated defaults.

Get API Key

Same as Mistral — uses `console.mistral.ai` keys.

Create Profile

```
cc-manage add
```

1. Select **Mistral Vibe** from the provider list
2. Enter filename (e.g., `mistral-vibe`)
3. Add new API key (or reuse Mistral key)
4. Accept defaults
5. Fetch live models or use defaults

Default Models

- `mistral-vibe-cli-latest`
- `mistral-medium-3.5`
- `devstral-small-latest`

Notes

- Mistral Vibe shares the same API endpoint as Mistral (<https://api.mistral.ai/v1>)
- Uses `MISTRAL_VIBE_API_KEY` by default but can fall back to `MISTRAL_API_KEY`

Launch

```
cc-switch mistral-vibe  
ccs
```

7. Codestral

Mistral's code-specialized model.

Get API Key

Same as Mistral — uses `console.mistral.ai` keys.

Create Profile

```
cc-manage add
```

1. Select **Codestral** from the provider list
2. Enter filename (e.g., `codestral`)
3. Add new API key
4. Accept defaults

Default Models

- `codestral-latest`
- `codestral-2508`

Notes

- Codestral chat goes through `https://codestral.mistral.ai/v1`
- FIM (fill-in-the-middle) is available upstream at `/v1/fim/completions` but `cc-manage` uses the chat endpoint

Launch

```
cc-switch codestral  
ccs
```

8. OpenCode Nemotron

NVIDIA's Nemotron models through OpenCode's API.

Get Auth Token

1. Go to opencode.ai
2. Sign up / log in
3. Get your API token from the dashboard

Create Profile

```
cc-manage add
```

1. Select **OpenCode Nemotron** from the provider list
2. Enter filename (e.g., `opencode-nemotron`)
3. Add new API key (your OpenCode auth token)

4. Accept defaults

Default Model

- nemotron-3-ultra-free

How It Works

- The proxy strips Claude-only metadata (thinking, context_management, output_config, etc.)
- Preserves tool calls through OpenAI function-calling fields
- Synthesizes Anthropic-compatible SSE from non-streaming upstream response
- Forces stream: false upstream regardless of client request

Launch

```
cc-switch opencode-nemotron  
ccs
```

9. NVIDIA NIM

NVIDIA's hosted models with dynamic discovery.

Get API Key

1. Go to build.nvidia.com
2. Sign up / log in
3. Navigate to **API Keys**
4. Generate a new key (starts with nvapi-)

Create Profile

```
cc-manage add
```

1. Select **NVIDIA NIM** from the provider list
2. Enter filename (e.g., nvidia-nim)
3. Add new API key
4. Fetch live models or use defaults

Default Models

- nvidia/nemotron-3-super-120b-a12b
- qwen/qwen3.5-397b-a17b
- qwen/qwen3.5-122b-a10b

Launch

```
cc-switch nvidia-nim  
ccs
```

10. Hugging Face

Access Hugging Face models through local proxy.

Get API Key

1. Go to huggingface.co/settings/tokens
2. Click **Create new token**
3. Copy the token (starts with hf_)

Create Profile

```
cc-manage add
```

1. Select **Hugging Face** from the provider list
2. Enter filename (e.g., huggingface-kimi)
3. Add new API key
4. Accept defaults

Default Models

- moonshotai/Kimi-K2.6
- moonshotai/Kimi-K2.5

Notes

- Kimi-K2.6 uses `chat_template_kwargs: {thinking: false}` for compatibility
- Kimi-K2.5 uses `thinking: {type: "disabled"}`

Launch

```
cc-switch huggingface-kimi  
ccs
```

11. DeepSeek

Direct Anthropic-compatible API — no proxy.

Get API Key

1. Go to platform.deepseek.com
2. Navigate to **API Keys**
3. Create a new key

Create Profile

```
cc-manage add
```

1. Select **DeepSeek** from the provider list
2. Enter filename (e.g., deepseek-chat)
3. Add new API key
4. Accept base URL: <https://api.deepseek.com/anthropic>

Default Models

- deepseek-chat
- deepseek-reasoner

Launch

```
cc-switch deepseek-chat  
ccs
```

12. Fireworks

Direct Anthropic-compatible API — no proxy.

Get API Key

1. Go to fireworks.ai
2. Navigate to **API Keys**
3. Create a new key

Create Profile

```
cc-manage add
```

1. Select **Fireworks** from the provider list
2. Enter filename
3. Add new API key
4. Accept base URL: <https://api.fireworks.ai/inference/v1/anthropic>

Default Model

- accounts/fireworks/models/llama-v3p3-70b-instruct

Launch

```
cc-switch <your-fireworks-profile>  
ccs
```

13. Together

OpenAI-compatible API through local proxy.

Get API Key

1. Go to api.together.xyz/settings/api-keys
2. Create a new key

Create Profile

```
cc-manage add
```

1. Select **Together** from the provider list
2. Enter filename
3. Add new API key
4. Accept proxy port (default: 18100)

Default Model

- meta-llama/Llama-3.3-70B-Instruct-Turbo

Launch

```
cc-switch <your-together-profile>  
ccs
```

14. xAI

OpenAI-compatible API through local proxy.

Get API Key

1. Go to console.x.ai
2. Navigate to **API Keys**
3. Create a new key

Create Profile

```
cc-manage add
```

1. Select **xAI** from the provider list
2. Enter filename
3. Add new API key
4. Accept proxy port (default: 18100)

Default Models

- grok-4
- grok-3

Launch

```
cc-switch <your-xai-profile>  
ccs
```

15. Ollama Cloud

OpenAI-compatible API through local proxy.

Get API Key

1. Go to ollama.com
2. Sign up / log in
3. Get your API key

Create Profile

```
cc-manage add
```

1. Select **Ollama Cloud** from the provider list
2. Enter filename
3. Add new API key
4. Accept proxy port (default: 18100)

Default Models

- gpt-oss:120b
- llama3.3:70b

Launch

```
cc-switch <your-ollama-profile>
ccs
```

16. Custom OpenAI-Compatible

Bring your own endpoint.

Create Profile

```
cc-manage add
```

1. Select **Any OpenAI-compatible cloud endpoint** from the provider list
2. Enter filename
3. Add your API key
4. Enter the upstream base URL (e.g., `https://your-provider.com/v1`)
5. Enter the proxy port (default: 18100)
6. Enter your model name(s)

Requirements

The endpoint must support: - POST /chat/completions (OpenAI Chat Completions format) - Authorization: Bearer <key> header - Standard messages, model, max_tokens fields

Launch

```
cc-switch <your-custom-profile>
ccs
```

Provider Comparison

Response Format Support

Provider	Text	Images	Tools	Streaming	Max Output
Anthropic	[x]	[x]	[x]	[x]	Model-dependent
Gemini	[x]	[x]	[x]	[x]	8192
OpenRouter	[x]	[x]	[x]	[x]	Model-dependent
Groq	[x]	[x]	[x]	[x]	4096
Mistral	[x]	[x]	[x]	[x]	Model-dependent
Mistral Vibe	[x]	[x]	[x]	[x]	Model-dependent
Codestral	[x]	[]	[x]	[x]	Model-dependent
OpenCode Nemotron	[x]	[]	[x]	[x] (fake)	Model-dependent
NVIDIA NIM	[x]	[x]	[x]	[x]	Model-dependent

Provider	Text	Images	Tools	Streaming	Max Output
Hugging Face	[x]	[x]	[x]	[x]	Model-dependent
DeepSeek	[x]	[x]	[x]	[x]	Model-dependent
Fireworks	[x]	[x]	[x]	[x]	Model-dependent
Together	[x]	[x]	[x]	[x]	Model-dependent
xAI	[x]	[x]	[x]	[x]	Model-dependent
Ollama Cloud	[x]	[x]	[x]	[x]	Model-dependent
Custom	[x]	[x]	[x]	[x]	Endpoint-dependent

Speed Tiers

Tier	Providers	Notes
Fastest	Groq, Anthropic	Sub-second response times
Fast	OpenRouter, DeepSeek, xAI	Low latency
Medium	Gemini, Mistral, Together, Fireworks	Good balance
Variable	NVIDIA NIM, Hugging Face, Ollama Cloud	Depends on model/load